


CARD DISPUTES AT A LATAM BANK

A disputed charge waits **37 hours** for a first answer.

The customer calls, a complaint is opened, and then nothing moves for a day and a half. When someone finally picks it up, the record doesn't say what the customer said.

**8,199**

dispute complaints in one year

69.8 %

still open or in process

n = 8,199

Recibra

Proof your claim is on record.

An AI service desk for one job, done end to end: answer account questions and take card dispute claims, in Spanish and Portuguese. Every conversation ends in one of three ways.

Resolves it

Finds the charge, checks the dispute rules, files the claim and reads it back before saying it's done.

Registré tu disputa. **DSP-1791...D5**

Asks or declines

When the request is unclear it asks one question. When it's outside scope, it says so instead of guessing.

¿Quieres disputar un cargo o ver tus cuentas?

Hands off, case ready

Fraud, legal threats or a request for a person go to staff with the customer's words, verified facts and open questions.

Te paso con un especialista. Prioridad: crítica.

LATAM Bank
Asistente · en línea

Me cobraron dos veces en Tienda Don José el 9 de mayo. Quiero disputar el cargo duplicado.

Resumen de la disputa

Transacción Tienda Don José · 9 may
Monto COP 199.435,05
Motivo cobro duplicado

Confirmar y enviar

Confirmar y enviar

Se ha presentado una disputa por un cargo duplicado de 199.435,05 COP.

✓ DSP-17912585753185C7152F9 · registrada

Escribe tu mensaje...


ONE DISPUTE, START TO FINISH

The customer sees a receipt. Staff see why.

DECISION TRACE · TURNS 1-2 · LIVE 2026-10-06

DISPUTE-POLICY.V1 · THRESHOLDS.V2

load_context understand route resolve_transaction check_eligibility · 5 rules pass confirm token ✓ file_dispute
verify · read-back ✓ reply



1 · Understand

The message is read and checked against thresholds.

2 · Confirm

The button sends the card's token; code checks it.

3 · Prove

Written, read back, then the receipt is shown.




Watch the full demo

youtu.be/YuovlUNRyoU


From 37 hours to a verified case number in seconds.

Today
to first response



37.0 h

Recibra
to a case number



9.1 s

Cost

model spend, priced models only (a lower bound)

≈ \$0.69
per 1,000 cases

2 / 358

unsafe outcomes (0.6 %)

63 / 65

handoffs with the right reason

0 / 65

missed transfers to a person

178 / 179

Portuguese cases answered correctly in Portuguese

What staff get today

one real complaint row

subcategory	Cargo no reconocido
description	"Queja relacionada con transactions"
conversation link	empty, in 67,095 of 67,095 rows
priority	Medium, no reason given

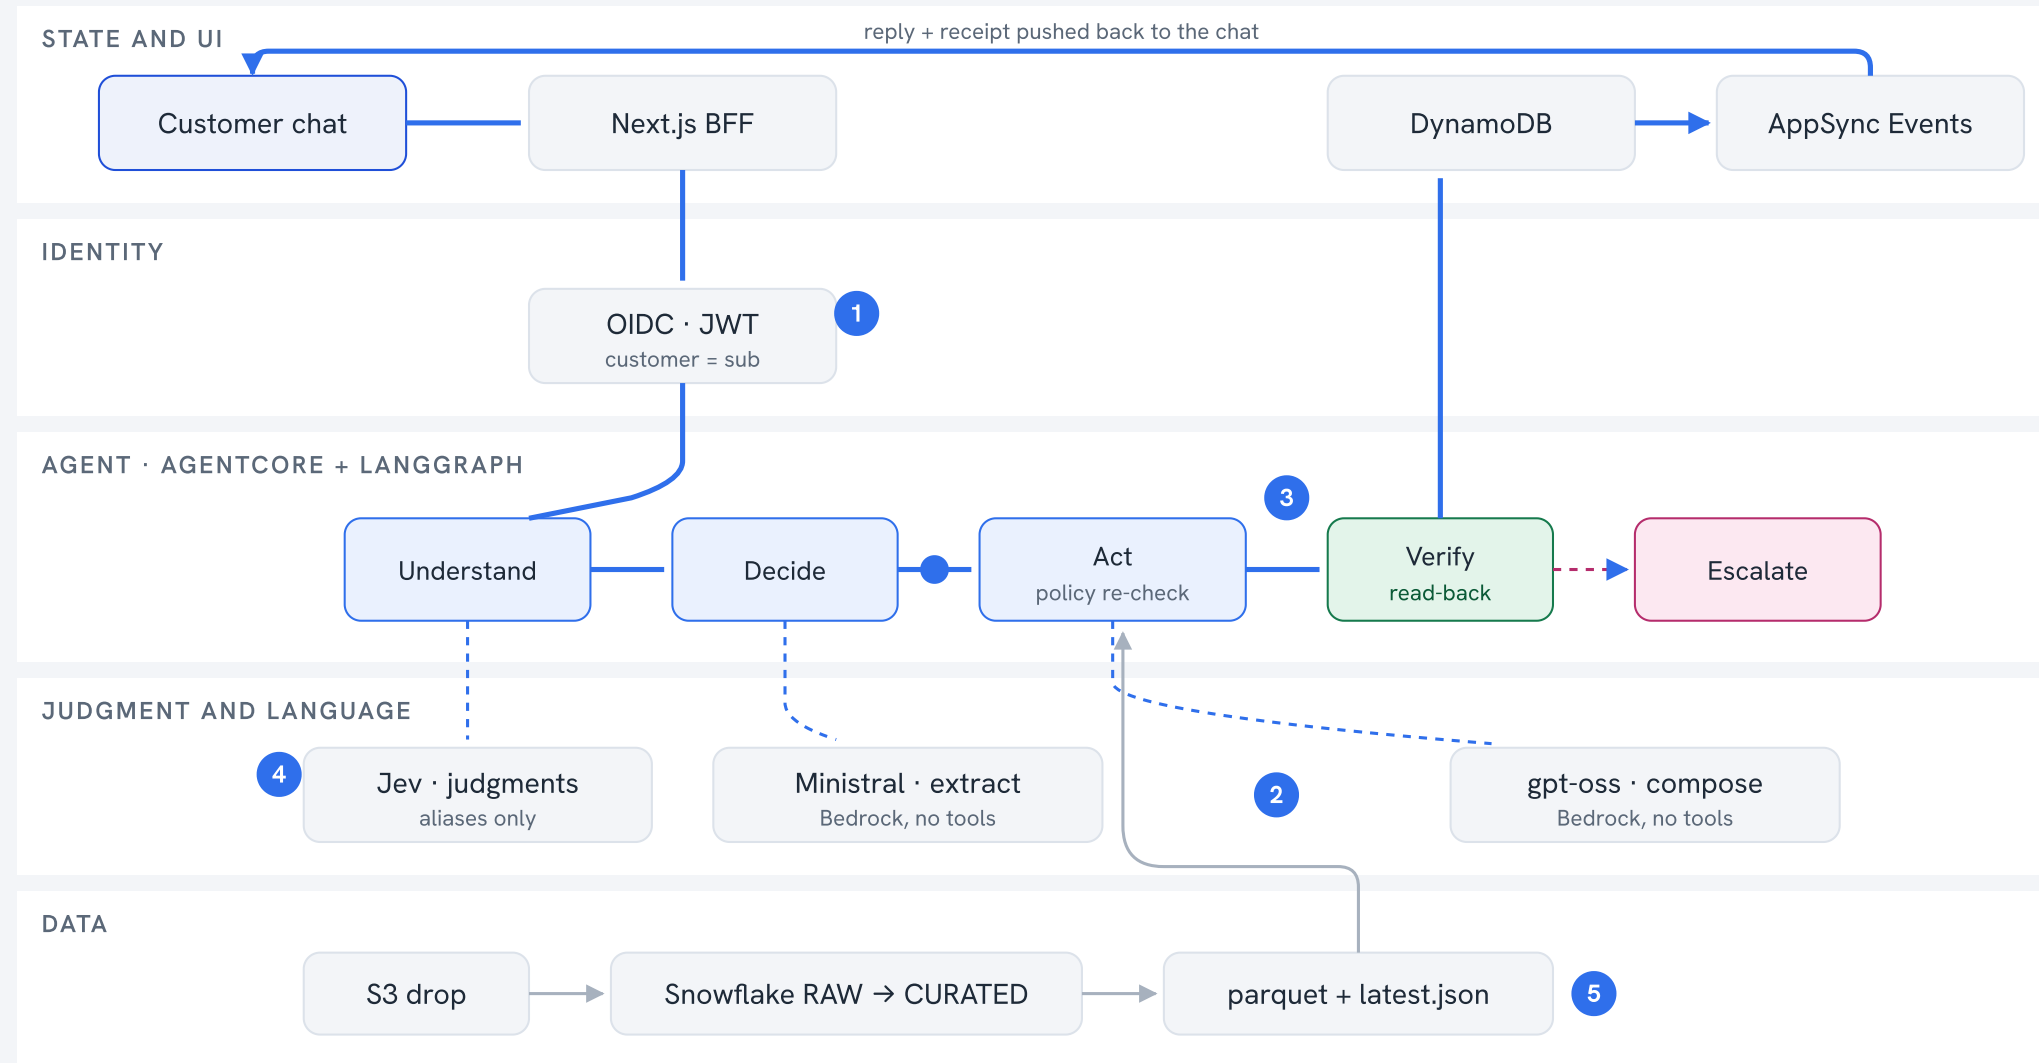
What Recibra hands over

handoff packet, live 2026-10-06

customer_request	"No reconozco un cargo en mi tarjeta, yo no hice esa compra"
why it's here	unauthorized use 0.97 vs 0.50 threshold → critical
open_questions	date and amount of the charge
transcript_ref	links to the chat; staff can take over

ONE DISPUTE TURN, END TO END

Built so the AI can't act on its own say-so.



- 1 Identity comes from the login token**
No one can read another customer's data by asking for it.
- 2 The language models have no tools**
Policy is re-checked inside the write itself, so persuasion can't grant permission.
- 3 Conditional write, then read-back**
We only tell the customer about actions that really happened.
- 4 No customer data leaves the bank**
Models never see customer or product ids or fraud signals; the judgment model picks transactions by alias.
- 5 7 Terraform roots, OIDC CI**
Reproducible from scratch, with no stored cloud keys.

LEARNED COMPONENT

On 150 blind test messages, the learned ranker acted on its own in 43 % of cases vs 28 % for the baseline, with 0 wrong actions. No proven gain overall.

From a working desk to a pilot

- 1** Real identity provider
Replace the mock login with the bank's own sign-in, with a second factor for staff.
- 2** Pilot against the 37-hour baseline
Measure real disputes, real customers, one country first.
- 3** Find the breaking point, staff the pager
10 customers at once: 30/30 answered, p95 6.9 s. Next, push until it saturates.
- 4** Retention rules and least-privilege CI
Decide how long transcripts and packets are kept.

WHAT WE KNOW IS MISSING

- Synthetic data; some patterns are generator artifacts
- No Portuguese-speaking customers in the data
- Judgment thresholds set by hand. On 358 held-out chats, 3 of 4 signals separate cleanly; injection overlaps
- Portuguese replies not yet reviewed by a native speaker
- Mock identity provider
- Load tested at 10 customers at once; saturation point unknown

**Recibra. Every dispute answered in seconds,
with proof.**

We report what we measured.